

Inter-process communication (IPC)

CS350
Kartik Gopalan

Some simple forms of IPC

- Parent-child
 - Command-line arguments,
 - `wait(...)`, `waitpid(...)`
 - `exit(...)`
- Reading/modifying common files
 - Servers commonly use 'pid' file to determine other active servers.
- Signals
 - Event notification from one process to another

CS350/BU

Some more forms of IPC...

- Shared Memory
 - Common piece of read/write memory.
 - Needs synchronization for access
- Semaphores
 - Locking and event signaling mechanism between processes
- Pipes
 - Uni-directional (if used cleanly)
 - 'ps -aux | more'
 - Can be used bi-directionally with some synchronization effort
- Sockets
 - Bi-directional
 - Not just across the network, but also between processes.

CS350/BU

Pipes

Pipe Abstraction


- Write to one end, read from another

- `pipe()`



CS350/BU


Parent-child communication using pipe



Here's an example.

CS350/BU

'ps -elf | more'



Here's an example.

CS350/BU

Being careful with read()/write()

- ❑ `read(fds[0], buf, 6);`
 - ❑ Doesn't mean read will return with 6 bytes of data!
 - ❑ It could be less. Why?
- ❑ Some reasons
 - ❑ `read()` could reach end of input stream (EOF).
 - ❑ Other endpoint may abruptly close the connection
 - ❑ `read()` could return on a signal.
- ❑ So you **MUST** incorporate error handling with every I/O call (actually with any system call)

CS350/BU

Error handling...

- ❑ You **must**
 - ❑ First check the return value of **every** `read(...)/write(...)` system call.
 - ❑ Then either...
 - ❑ Wait to read/write more data OR
 - ❑ Handle any error conditions

```

More convenient to write a wrapper function
/* Write "n" bytes to a descriptor. */
ssize_t writen(int fd, const void *vptr, size_t n)
{
    size_t nleft;
    size_t nwritten;
    const char *ptr;

    ptr = vptr;
    nleft = n;
    while (nleft > 0) {
        if ((nwritten = write(fd, ptr, nleft)) <= 0) {
            if (errno == EINTR)
                nwritten = 0; /* call write() again */
            else return(-1); /* error */
        }
        nleft -= nwritten;
        ptr += nwritten;
    }
    return(n);
}
  
```

CS350/BU


Shared Memory, Semaphores

Man pages : `shmget`, `shmat`, `shmdt`, `shmctl`
`semget`, `semop`, `semctl`

Also in Section 15.8 and 15.9 in APUE book

Shared Memory

Common chunk of read/write memory
among processes



CS350/BU

Creating Shared Memory

```
int shmget(key_t key, size_t size, int shmflg);
```

Example:

```
key_t key;
int shmid;

key = ftok("<somefile>", 'A');

shmid = shmget(key, 1024, 0644 | IPC_CREAT);
```

Here's an example.

CS350/BU

Attach and Detach Shared Memory

```
void *shmat(int shmid, void *shmaddr, int shmflg);
int shmdt(void *shmaddr);
```

Example:

```
key_t key;
int shmid;
char *data;

key = ftok("<somefile>", 'A');
shmid = shmget(key, 1024, 0644);
data = shmat(shmid, (void *)0, 0);
...
shmdt(data);
```

Here's an example.

CS350/BU

Deleting Shared Memory

```
int shmctl(int shmid, int cmd, struct shmid_ds *buf);
```

```
shmctl(shmid, IPC_RMID, NULL);
```

Example

CS350/BU

Command-line IPC control

ipcs

- Lists all IPC objects owned by the user

ipcrm

- Removes specific IPC object

CS350/BU

Signals

Signals Overview

- Signal is a notification to a process that an event has occurred.
 - Could come from another process or from the OS
- Type of event determined by type of signal
- Try listing all signal types using

```
% kill -l
```
- Some interesting signals
 - SIGCHLD, SIGTERM, SIGKILL, SIGSTOP

CS350/BU

Handling Signals

- Signals can be caught - i.e. an action can be associated with them
 - SIGKILL and SIGSTOP cannot be caught.
- Actions to signals can be customized using `sigaction(...)` which associates a signal handler with the signal.
- Default action for most signals is to terminate the process
 - Except SIGCHLD and SIGURG are ignored by default.
- Unwanted signals can be ignored
 - Except SIGKILL or SIGSTOP
- Here's an example

CS350/BU

More on SIGCHLD

- ❑ Sent to parent when a child process **terminates** or **stops**.
- ❑ If `act.sa_handler` is `SIG_IGN`
 - ❑ `SIGCHLD` will be ignored (default behavior)
- ❑ If `act.sa_flags` is `SA_NOCLDSTOP`
 - ❑ `SIGCHLD` won't be generated when children stop
- ❑ `act.sa_flags` is `SA_NOCLDWAIT`
 - ❑ children of the calling process will not be transformed into zombies when they terminate.
- ❑ These need to be set in `sigaction()` *before* parent calls `fork()`

CS350/BU

How to avoid zombies?

- ❑ Parent could install a signal handler for `SIGCHLD`
- ❑ Call `wait(...)/waitpid(...)` inside the signal handler

```
void handle_sigchld(int signo) {  
    pid_t pid;  
    int stat;  
  
    pid = wait(&stat);  
    printf("child %d terminated\n", pid);  
}
```

- ❑ Here's an example.

CS350/BU

More information...

- ❑ Check '`man sigaction(...)`'
- ❑ Understand what happens when signal is delivered in the middle of a system call?
 - ❑ Different OSes have different behavior.
- ❑ Google for keywords "Unix Signals"
 - ❑ Tons of useful links

CS350/BU

References

- ❑ Unix man pages
- ❑ "Advanced Programming in Unix Environment" by Richard Stevens
 - ❑ <http://www.kohala.com/start/apue.html>

CS350/BU