Live Migration of Virtual Machines

Christopher Clarke, Keir Fraser, et. al.

NSDI 2005

What is live migration?

- Move a VM from one physical machine to another even as its applications continue to execute during migration
- Live VM migration usually involves
  - Migrating memory state
  - Migrating CPU state
  - Optionally, migrating virtual disk state

Why Migrate VMs Live?

- Load Balancing
- System Maintenance
- Avoiding residual dependencies at source host which occurs with process migration
  - E.g. system call redirection, shared memory
- Avoiding Lost Connections

Performance Goals in Live Migration

- Minimizing Downtime
- Reducing total migration time
- Avoiding interference with normal system activity
- Minimizing network activity

Migrating Memory

- Pure stop-and-copy
  - Freeze VM at source,
  - Copy the VM's pseudo-physical memory contents to target,
  - Restart VM at target
  - Long downtime.
  - Minimal total migration time = downtime
- Pure Demand Paging:
  - Freeze VM at source,
  - Copy minimal execution context to target
    - PC, Registers, non-pageable memory
  - Restart VM at target,
  - Pull memory contents form source as and when needed
  - Smaller downtime
  - Sloooow warm-up phase at target during page-faults across network

Pre-copy migration

- DON'T freeze VM at source ➔Let it continue to run
- Copy VM's pseudo-physical memory contents to target over multiple iterations
  - First iteration ➔ copy all pages.
  - Each subsequent iteration ➔ copy pages that were dirtied by the VM during the previous iteration
- Xend – a daemon in Domain 0 – maps the guest VM's address space and transfers the pages over TCP connection to the target.
- Do a short stop-and-copy when number of dirty pages is "small enough".
- But what if number of dirty pages never converges to a small enough number?
  - After a fixed number of iterations, give up and stop-and-copy.

Stages of Migration

1. Pre-Migration
   - Prepare the guest VM for migration via event channel notification
2. Reservation at target
   - Check if target has enough resources to receive the migrating VM
3. Iterative Pre-Copy
   - Copy memory contents over multiple rounds
4. Stop-and-Copy (downtime)
   - Freeze the guest and copy any residual state, including remaining dirty memory pages.
5. Commitment
   - Indicate to target machine that all state has been transfered
6. Activation
   - Target m/c restarts the guest

So what's the catch? How do we track dirtied pages?

- Mark the VM's memory pages as read-only after each iteration.
- Trap write operations via hypervisor to xend and track dirtied pages.
- Reset after each iteration
- Works well as long as writes are infrequent

Managed Migration

- Migration initiated and managed outside of the VM
  - Typically by xend running in Dom0 at both the source and target
- Xend at source contacts xend at target and transfers the VM state across the network.

Self Migration

- Guest OS migrates itself (mostly)
- Xend on source machine not involved.
- Migration stub needed at destination
- Challenge:
  - OS must continue to execute while transferring its final state.
  - Perform a careful (complicated) 2-stage checkpoint and copy.

Minimizing impact on running services

- Dynamically adapt bandwidth limit.
- Use minimum bandwidth in first round
- Calculate dirtying rate in each subsequent round to determine bandwidth

Other tricks

- Stun Rogue Processes
  - Those that don't stop dirtying memory
- Free Page Cache Pages
  - Can be re-cached at target
  - Potential performance hit

Migrating Network Connections

- Migrating VM carries its
  - IP address,
  - MAC address, and
  - all protocol state, including any open sockets
- So nothing special to do while migrating within a switched LAN environment.
- What about the backward (re)learning delay at the network switches?
  - Switches needs to re-learn the new location of migrated VM's MAC address
  - Solution: Send an unsolicited ARP reply from the target host.
  - Intermediate switches will re-learn automatically.
  - Few in-flight packets might get lost.

Storage Migration

- Much bigger problem
  - Many gigabytes of local disk image possible.
- Bypass the problem
  - Assume the storage is over the network and remains accessible from the new target machine.
  - E.g. Network File System (NFS), or Network Block Device(NBD), or iSCSI etc.