

Low Bandwidth VM Migration Via Opportunistic Replay

Context of the paper

- ❑ Migration of VMs used in mobile computing
- ❑ Non-live migration
 - Suspend VM at the source host
 - Transfer state to destination
 - Resume VM at destination

Problem

- ❑ VMs are large
 - Often tens of GB in size
- ❑ Mobile devices have low bandwidth
- ❑ Implies, excessive transfer time
- ❑ Goal: Reduce transfer time by reducing amount of information transferred.

Solution

- ❑ Opportunistic replay
 - Optimize migration between frequently visited hosts
- ❑ Key observation: Replay does not have to be perfect in order to be useful

Opportunistic replay

- ❑ Log user interactions with a VM, such as keystrokes, mouse clicks etc.
- ❑ Ship the log to the destination
- ❑ Replay the log on an identically configured VM which has the same initial state.

The catch?

- ❑ User interactions are not the only factors that affect the VM state.
 - Hardware interrupts, network packet processing, background system tasks, etc. all effect the final VM state.
- ❑ So replay will produce state that is close, but not identical to the desired final state.
- ❑ Solution: After replay, find out and ship only the differences in state between the source and the target.
- ❑ Worst-case: difference is large and we ship large amount of state.
- ❑ But correctness is never lost.

Challenge 1: Incomplete log capture

- ❑ Capturing all external stimuli requires
 - External interrupts
 - Data transfers
 - User input
- ❑ Can result in a very large log
 - Observed up to 1.2GB per day
- ❑ Also capturing all external events not possible in closed-source VMMs such as VMWare
- ❑ Solution: Just capture user-interactions
 - Most windowing systems provide interface to interpose logging code.
 - Will produce short, but incomplete, logs

Challenge 2: Non-deterministic Events

- ❑ Network access to websites with dynamic content
 - Replay of user click may return different content at the target.
 - E.g. Stock quotes, Parasitic content such as ads.
 - In most cases, users don't care about identical content.
- ❑ How about interrupts?
 - Ok, as long as it doesn't impact correctness or quality of user experience.

Challenge 3: Exactly once side effect

- ❑ Some events should not be replayed.
 - E.g. Stock purchase, online bank payments, sending emails.
- ❑ Solution 1: Block outbound network messages during replay.
- ❑ Solution 2: Detect exactly-once actions and skip them during replay. How?

Prototype implementation

- ❑ Xen 3.1
- ❑ Custom block driver that exports a virtual disk to VM.
 - Keeps track of dirty disk blocks and dirtying order
- ❑ rsync to synchronize (?) VM memory images.
- ❑ Xnee: X11 tool for record-replay.
- ❑ Doesn't address exactly-once problem.

Prototype (contd.)

- ❑ Initially both source and destination have identical copies of suspended VM.
- ❑ Replay speed may be faster than capture speed
 - Suppress the user think time.
- ❑ Modified disk chunks shipped to destination in the background while replay is in progress.
 - In reverse order (?)
- ❑ Differences in final state at source and destination computed
 - SHA-1 Hash
 - Applied on replayed image at target