

Transparent Network Services via a Virtual Traffic Layer for Virtual Machines

Motivation

- HPC community more interested in using VMs
- Particularly because of
 - VM migration capabilities
 - Much more cleaner than process migration
 - Transparent addition of new services without changing the VM or the applications.
- Monitoring and controlling
 - Execution of VMs
 - Network communication of VMs

Problem being addressed

- How to provide new class of network services to applications in the VMs?
 - Without modifying either the VM or the Apps.
- Goals:
 - Monitor traffic
 - Control routing
 - Interpose and modify data and signaling

Contribution

- VTL: Virtual traffic later toolset
 - Packet capture
 - Packet inspection
 - Connection state maintenance and modification
 - And combination of the above
- Basically a fancy NAT for VMs

Service Examples

- Tor-VTL
- Subnet Tunneling
- Local Acks
- Split TCP
- Protocol transformations
- Stateful firewall
- TCP Keep-alives
- Traffic wormholing for IDS

VTL Toolset

- Library of packet monitoring and manipulation functions
- User-level code that relies on libpcap or winpcap
- Assumption: All traffic goes through a virtual network interface (“host-only” NIC) via domain 0 or equivalent.

Relationship to VNET

- VNET
 - Allows nodes on disjoint IP subnets to pretend as if they are on the same subnet.
 - Relies on a set of proxying hosts that perform MAC address transformations in packets while keeping IP addresses the same.
- VTL can work either standalone or with VNET
 - VTL-VNET communication via a local channel.

VTL Example

```
RawEthernetPacket pkt;
iface_t * src_if = if_connect("src_device");
iface_t * dst_if = if_connect("dst_device");

while (if_read(src_if, &pkt)) {
    if_write(dst_if, &pkt);
}
```

Figure 1: Simple one-way VTL bridge.

```
RawEthernetPacket pkt;
unsigned long dst, new_dst;

dst = *(uint32 *)IP_DST(pkt.data);
*(uint32 *)IP_DST(pkt.data) = new_dst;

dst = GET_IP_DST(&pkt);
SET_IP_DST(&pkt, new_dst);
```

Figure 3: Example of basic packet access.

```
int create_data_pkt(vtl_model_t * model,
                  char * data,
                  int data_len) {
    RawEthernetPacket data_pkt;
    create_empty_pkt(&model,
                    &data_pkt,
                    OUTBOUND_PKT);
    memcpy(TCP_DATA(data_pkt.data),
          data, data_len);
    ip_len = GET_IP_TOTAL_LEN(data_pkt.data);
    ip_len += data_len;
    SET_IP_TOTAL_LEN(data_pkt.data, ip_len);
    compute_ip_checksum(&data_pkt);
    compute_tcp_checksum(&data_pkt);
    sync_model(&model, &data_pkt);
    pkt_len = data_pkt.get_size() + data_len;
    data_pkt.set_size(pkt_len);
    queue_pkt(&data_pkt);
    return 0;
}
```

Figure 4: Creating a data packet.

User space overhead




Figure 2: Performance overhead of VTL.

State Model

- Basically connection state information for individual network connections in VM
- For TCP
 - Sequence and ACK numbers
 - Timestamps
 - Receive window size?
 - IP state (?) etc
- Initialization
 - Manually
 - Supplying an example packet

Anonymous TOR Service




Figure 5: Network configuration of Tor-VTL.

- Use of SOCKS standard for proxies
- Four states:
 - Open, Established, close, Error
- DNS Lookups
 - SOCKS4a
- ARP interception

Subnet Tunneling

- Optimize communication between two VM on the same LAN but different VNET sIP subnets
- Basically MAC address remapping


Figure 6: Subnet tunneling example.

Figure 6: Subnet tunneling example.

Enhancing Network Performance

- For high bandwidth-delay product networks
- Local ACKs
- Split TCP
- Protocol transformations

Latency (ms)	TCP (Mbps)	TCP-ROCKSUF (Mbps)	UDT (Mbps)	Other (Mbps)
1	80	40	0.1	0.1
10	15	15	15	0.1
100	1.5	1.5	15	0.1
500	0.3	0.3	12	0.1
1000	0.15	0.15	12	0.1



Connection persistence during VM migration

- Over wide-area, we have longer VM migration time during which VM does not run.
- Routing changes: Handled by VNET
- Timeouts: handled by VTL
- TCP keep-alive packets
- Advertising a receive window size of zero to peer
- Respond to periodic probes from peer

- Over wide-area, we have longer VM migration time during which VM does not run.
- Routing changes: Handled by VNET
- Timeouts: handled by VTL
- TCP keep-alive packets
- Advertising a receive window size of zero to peer
- Respond to periodic probes from peer

Cooperative wormhole scheduling and Vortex

The diagram illustrates the system architecture. On the left, a 'Commodity PC' contains a 'Windows/UNIX' layer, an 'Apps' layer, an 'Operating System' layer, a 'Vortex' component, a 'VTL' layer, a 'Firewall' layer, an 'iSCSI' layer, and an 'NIC' layer. The 'Vortex' component is connected to a 'VNET Proxy' box. The 'VNET Proxy' is connected to a 'VNET Overlay' cloud. The 'VNET Overlay' is connected to a 'VM' box labeled 'VM Based HoneyPot'. The 'VM' is connected to an 'IDR Analysis BackEnd' box. The 'IDR Analysis BackEnd' is connected to a 'Physical HoneyPot' box. The 'Physical HoneyPot' is connected to a 'Switch' box, which is part of the 'Backend Network'.

