

Wireless Ad Hoc Routing Without Explicit Control Messages

Kartik Gopalan Navodaya Garepalli Ping Yang
Computer Science, State University of New York at Binghamton

Abstract—This paper introduces a new wireless ad hoc routing protocol called *Ad hoc Bridging Protocol* (ABP). ABP distinguishes itself from earlier proposals in that it uses no explicit control messages for route discovery, setup, and maintenance, while making minimal use of implicit (data-like) control messages that require no special processing at intermediate nodes. ABP is inspired from the self-learning plug-and-play nature of traditional transparent bridges in wired LANs. Like wired bridges and many earlier ad hoc protocols, nodes in ABP learn and adapt routes to different destinations on-the-fly by observing the packets sent by those destinations. As a result, ABP is completely decentralized, self-starting, and automatically adapts to varying network conditions. However, a major difference is that ABP does not attempt to prevent routing loops at all costs. Consequently, ABP does not construct or maintain any overarching control infrastructure for loop prevention, such as extensive network-wide spanning trees in traditional LAN bridges, destination sequence numbers in AODV, or source-routing in DSR. The essential idea is to couple the backward learning technique from transparent bridging, with an associated refresh mechanism, and the use of an identification field to effectively limit the impact of any transient loops. We also prove that transient loops, if and when they occur, do not last longer than a bounded, short time interval and do not negatively impact the network performance. Results demonstrate that, even without explicit control messages, ABP can perform competitively in comparison to AODV and DSR protocols while significantly reducing the protocol complexity.

I. INTRODUCTION

State-of-the-art wireless ad hoc routing protocols rely upon the use of explicit (and possibly piggybacked) control messages with an associated extensive control infrastructure to set up and maintain loop-free routes. The use of explicit control messages enables these protocols to optimize for network throughput, such as by reducing broadcast messages. While the on-demand nature of these protocols and several optimizations can reduce the explicit control message overhead, the fact remains that a large fraction of protocol complexity itself arises from the need to maintain this control infrastructure and process the explicit control information.

In this paper, we propose a new wireless ad hoc routing protocol, called *Ad hoc Bridging Protocol* (ABP), a simple and as-yet-overlooked mechanism to perform ad hoc routing that greatly reduces protocol complexity. ABP neither uses any *explicit* control messages nor maintains any overarching control infrastructure for route set up, maintenance, or loop prevention. The key idea is to combine the on-demand nature of ad hoc routing protocols with the simplicity of *backward learning* bridges. While the basic principle of backward learning itself is not new, ABP employs this principle without

having to construct and maintain network-wide spanning trees as in wired LANs, or use destination sequence numbers as in AODV [22], or resort to source routing as in DSR [12]. Our results show that ABP can deliver highly competitive network performance, if not better, in comparison to the popular AODV and DSR protocols.

Just as in wired bridging and many wireless ad hoc protocols, nodes in ABP use backward learning to set up and adapt routes to different destinations on-the-fly. However, unlike conventional wireless ad hoc protocols, ABP nodes implicitly learn routes from data packets themselves, rather than through explicit control messages. In the case of unidirectional traffic flow, ABP selectively uses minimal number of *implicit* (or data-like) control messages that do not require any special processing in the network interior. At the same time, unlike wired bridges, ABP also does not go out of its way to prevent routing loops at all cost but limits their performance impact by design, if and when such loops arise. Hence ABP eliminates the overhead of constructing and maintaining an extensive infrastructure for loop prevention, such as network-wide spanning trees in wired-LAN bridges, destination sequence numbers in AODV, or source-routing in DSR. We prove that any transient loops, which may potentially arise, do not last beyond a bounded small duration. Moreover, packets do not traverse around transient loops more than once and thus do not lead to exponential packet storms.

Additional features of ABP include no need for promiscuous mode operation of wireless interfaces, no *Hello* messages between neighbors, and elimination of ARP (Address Resolution Protocol) request/reply messages. Thus, ABP is completely decentralized, self-starting, and automatically adapts to the widely varying network conditions. Security issues in ABP are fundamentally the same as in any protocol that relies upon the backward learning principle, including AODV and wired bridges. These issues are currently being investigated by our group, even though they are not addressed in this paper. We could also, in principle, adapt the concepts proposed for ABP to an existing protocol such as AODV to eliminate control message processing. However, to maintain clarity of exposition, we choose to describe these concepts in the context of a new ABP protocol and leave their adaptation to other protocols for future research.

II. NETWORK MODEL

We consider a collection of possibly mobile nodes operating together in an ad hoc network without any infrastructure

support. ABP nodes have two levels of identity. First is the link level (or layer-2) address which is visible only to immediate neighbors of the node. Second is the network level (or layer-3) address which a node uses to identify and establish communication sessions with any other node (not necessarily a neighbor). Without loss of generality, in this paper we assume that layer-2 address is the 6-byte 802.11 MAC address and the layer-3 address is the 4-byte IPv4 address (or IP address). ABP nodes do not assume any form of hierarchical or structured relationship among the different layer-3 or layer-2 addresses of nodes in the ad hoc network. All nodes in the network can participate in ABP. A node does not need to have any prior knowledge of the identity, connectivity, or location, of other nodes in the network except, of course, the layer-3 IP address of the nodes with which it explicitly wishes to communicate. A node dynamically learns and unlearns the layer-2 MAC address identity of its neighbors only when it needs to communicate with/via those nodes. Nodes do not exchange any periodic Hello messages to maintain neighbor information. The backward learning mechanism assumes a bidirectional ability to communicate among immediate neighbors. Hence ABP operates on top of link layer protocols, such as the 802.11 family of protocols, which support link-level acknowledgment mechanism for packet transmission between neighbors. Link-level acknowledgments help ABP nodes to detect and avoid unidirectional links, which may occur at times due to differing antenna or propagation patterns. The only capability required from the network interface hardware is the ability to send/receive unicast packets (with link-level acknowledgments) and layer-2 broadcast packets (without link-level acknowledgments). A network interface is not required to operate in ‘promiscuous’ mode (which allows an interface to receive every transmission in the wireless neighborhood without layer-2 address filtering).

III. ABP OVERVIEW

ABP is a completely on-demand ad hoc routing protocol, in which nodes dynamically learn routes by applying the concept of *backward learning* from data packets in transit. One important consequence of not using any network-wide control infrastructure or explicit control messages is that the ABP operation at each node is fully independent. Every node makes completely local routing decisions that require no coordination with other nodes. ABP is positioned as a thin protocol layer in the networking stack between the Layer-3 (IP layer) and the Layer-2 (MAC layer). Although forwarding is based on destination IP address, the ABP routing is slightly different from traditional IP routing. As we will detail in Section III, the next-hop address for a destination IP in the routing table is the layer-2 MAC address of the next-hop rather than its IP address. ABP nodes do not care about the IP-to-MAC address mappings of other nodes, except as needed to make forwarding decisions.

In the absence of any control infrastructure for loop prevention, ABP does not try to avoid loop formation at all times. Rather the ABP nodes utilize the Identification and time-to-live




Fig. 1. Illustration of route discovery mechanism.

(TTL) fields carried within each packet’s IP header to contain the impact of any potential loops that may arise. (In non-IP networks, packets that do not carry these pieces of embedded information can be handled by encapsulating them within a small header that carries a TTL and a unique identification value.) Even if transient routing loops emerge under some situations, we prove that they last for a bounded small duration. Similar mechanisms to get around the restrictions imposed by spanning tree in wired-LAN bridges has been advocated by Perlman [23].

Routing Table and Data Packets: Each ABP node maintains its forwarding knowledge in a *soft-state* routing table. Each routing table entry contains the following basic information: (1) Destination IP Address (D), (2) Next Hop MAC Address (m), (3) A list of alternative next hop MAC addresses, (4) Valid flag (v), and (5) Expiration time (e). Only routing entries that are *valid* can be used for packet forwarding. Routing entries can be invalid during the times when a node is recovering after a failure or mobility of a neighbor.

A routing entry gets *refreshed* each time a packet from the corresponding destination D is received from the next-hop MAC address m . In other words, *reverse traffic from the destination D updates its forward routing entries at intermediate nodes*. This is unlike in the case of AODV, where forward traffic towards D is used to refresh the forward routing entries. Expiration Time field indicates the amount of time left before an *un-refreshed* routing entry can be purged.

We represent the relevant header fields of a packet by the quadruple $[s_m, d_m, S_{IP}, D_{IP}]$. Here s_m and d_m represent the source and destination MAC addresses of immediate neighbors exchanging the packet and S_{IP} and D_{IP} represent the original source and final destination IP addresses. A $*$ represents a broadcast MAC or IP address.

IV. ROUTE DISCOVERY

The route discovery process in ABP has similarities to that in AODV protocol which also employs backward learning. However, AODV performs learning using explicit control messages, such as RREQ and RREP, rather than implicitly through data packets. Figure 1 illustrates the basic route setup mechanism in ABP. Source S initiates communication with destination D for the first time by flooding its first data packet with header $[s_m, *, S_{IP}, D_{IP}]$. This data packet travels towards D , setting up routing table entries for S at intermediate nodes. Any intermediate node X , that receives the packet from S , processes the packet in two phases: the *learning* phase and the *forwarding* phase.

Note here that there is an inherent tradeoff between flooding a first data packet (as in ABP) versus flooding a first control

packet (as in AODV and DSR). Flooding a (potentially large) data packet could be more expensive than flooding a (smaller) control packet. In practice, most communication sessions begin with an initial bidirectional exchange of small packets (such as TCP SYN/ACK packets) which greatly reduces the flooding cost during route setup. Further, ABP significantly limits the impact of any network-wide broadcasts via mechanisms described later in Section IV-IV-C and IV-IV-D. Additionally the elimination of explicit control messages and the associated loop prevention infrastructure greatly simplifies the protocol complexity.

A. Learning Phase

In the learning phase, node X learns about the direction in which node S is reachable. First consider the case in Figure 1(a) where X is an immediate neighbor of S , but is not aware of this fact and does not have any prior routing information for S in its routing table. When X receives the first broadcast packet with header $[S_m, *, S_{IP}, D_{IP}]$, it immediately infers that a node with layer-3 IP address S_{IP} can be reached in the direction of an immediate neighbor with layer-2 MAC address S_m . The node X then enters this newly learnt forwarding information in its routing table as a valid routing entry for S_{IP} . Note that X does not care that S_{IP} and S_m are the same neighbor's IP and MAC addresses respectively. All X cares is that S_{IP} is reachable in the direction of S_m .

Consider the second case in Figure 1(b) when X is not an immediate neighbor of S and does not have any prior routing entry for S_{IP} . Rather X receives a packet with header $[W_m, *, S_{IP}, D_{IP}]$, forwarded from another neighbor W , which does not know the next hop to D . Here X will learn that S_{IP} is reachable in the direction of W_m .

Finally consider the third case, again in Figure 1(b), when node X already has a valid routing table entry for S . Assume that a non-duplicate (i.e. first-time) packet with header $[W_m, *, S_{IP}, D_{IP}]$ is received from a neighbor W , which X already knows is a valid next hop for S_{IP} . In this case the only new information that node X can learn from the packet is that S_{IP} is still reachable along the direction of W_m and can refresh its routing table entry for S by resetting the expiration time field. On the other hand, if the current next hop for S_{IP} is not W_m then X can infer that either the earlier known route to S_{IP} has failed, or S has moved to a new location, or simply a better route to S may have appeared depending on the specific link/path quality metric of interest. In this case, X can choose to *re-learn* the new next-hop to S_{IP} , as described later in Section V.

Another function of the learning phase is to learn alternative routes. Suppose an intermediate node X receives multiple copies of the same packet, which originated from S , from different neighbors. The first copy of the packet will be used by X to learn the *primary* next hop towards S . Any subsequent copy of the packet will be recognized by X as a duplicate (using mechanism is described later in Section VI). Before dropping the duplicates, X will extract the source MAC

address from the duplicates as *alternative* next hop towards S ³. This alternative next hop can be used to replace the the primary next hop, in case the primary link fails or the alternative has better quality metric (such as higher bandwidth or lower noise) than the primary. While the consideration of alternative link/path quality metrics other than delay might be less of an issue in a mobile ad hoc environment in the sense that changes in the network can quickly render the quality information stale, such metrics might be important in mesh networks with largely static topology. Hence ABP permits the use of different quality metrics, whenever they are available.

One consequence of the backward learning mechanism described above is that ABP automatically populates the routing table with the mappings between destination IP address and next-hop MAC address. *Hence ABP eliminates the need to separately invoke the Address Resolution Protocol (ARP) to resolve these mappings*, which would otherwise introduce additional network traffic in the form of ARP Request/Reply messages. Another consequence is that neighbors need to learn of each others' existence only when they need to forward data packets. There is no particular need for nodes to proactively maintain neighbor information using Hello messages.

B. Forwarding Phase

In the forwarding phase, node X first checks if the packet is destined to itself, i.e. $D_{IP} = X_{IP}$. If so, the packet is delivered to the IP layer (layer-3) within X . If not, then X determines the next hop for destination D . If no valid routing entry exists for D_{IP} then, as shown in Figure 1(a), X re-broadcasts the data packet to its neighbors with a modified header $[X_m, *, S_{IP}, D_{IP}]$. If X already has a valid routing entry for D_{IP} with the next hop neighbor Y then, as shown in Figure 1(b), X forwards the data packet to Y with a modified header $[X_m, Y_m, S_{IP}, D_{IP}]$. *Hence, if parts of the network already have valid routes to D then those routes are reused.*

If, for some reason, the transmission to next hop node Y fails, then X invalidates the routing entry and attempts to send the data packet to one of its neighbors Z in the *alternative* next hop list of routing entry for D with a modified header $[X_m, Z_m, S_{IP}, D_{IP}]$. If transmissions to none of the alternative next hops succeed, then X simply re-broadcasts the packet with a modified header $[X_m, *, S_{IP}, D_{IP}]$.

To complete the picture, by the time the packet from S reaches D , the intermediate nodes in the network (including D) learn about the route to reach S . Similarly, when the receiving user-level application at node D responds back with its own data to its peer application at node S , the intermediate nodes (including S) learn about how to reach D as well.

C. Handling Silent Endpoints

One of the distinguishing features of ABP is the manner in which it handles connections that involve silent endpoints. Many real-world network applications engage in bidirectional communication where both sides exchange either data or control packets at application or transport level. For instance, the popular TCP protocol would, at the minimum, involve at

least one end transmitting data and the other end responding back with ACKs. In the unlikely event, where only one of the end-points actively sends traffic, the backward learning mechanism described above would maintain routing information for only the active sender and not for the passive receiver. This creates an undesirable situation where all data packets to the silent receiver might be flooded due to absence of routing information for the receiver.

To rectify this situation, the ABP layer at the receiver node D monitors whether packets from S are being received for some *ACTIVITY_INTERVAL* duration (say 80% of maximum route expiration time) without the application at D sending any data back to S . If so, the ABP layer at D transmits a **dummy IP packet** to S_{IP} with zero byte data payload. As the dummy IP packet from D_{IP} to S_{IP} travels through the network, intermediate nodes learn the routing information for D_{IP} . Thus future data packets from S to D are unicasted rather than flooded through the network. Since the dummy IP packet is not meant for any specific application at S , it is silently discarded by the IP layer at S . ABP layer at D stops sending dummy IP packets once it observes that S is not sending any more data packets.

The dummy IP packet is an *implicit control message* for the following two reasons. First, the dummy IP packet is treated just like any other data packet and forwarded by the intermediate nodes without any special processing. Secondly, this packet is created only under the exceptional circumstance when one of end-points is completely silent for a long period.

D. Quelling Back-to-back Broadcasts

When TCP is used for communication, the session would normally begin with a transport level 3-way handshake, which would help to set up routing information for both source S and destination D . However, for non-TCP sessions, the application initiating the communication at source S might, in some cases, begin by sending multiple back-to-back packets to the destination D . In case the intermediate nodes do not have routing information for D , all the initial back-to-back packets from S to D might be flooded through the network. This situation would persist until either the application at D responds with its own data, or the ABP layer at D sends a dummy IP packet to S . When this situation happens, it may result in undesirable back-to-back network-wide broadcasts for a small duration of time.

To mitigate this situation, we can apply an optimization at the ABP layer of the initiator S . After sending out the first packet, the ABP layer at S buffers the remaining back-to-back outgoing packets until a data packet is received from D . Since node D is guaranteed to transmit either an application data packet or a dummy IP packet within a bounded *ACTIVITY_INTERVAL*, the space set aside at S for this buffer can be bounded and small. After receiving a packet from D , S can be certain that routing state is set up in the network for both end-points and S can now safely send out the remaining buffered packets as well.

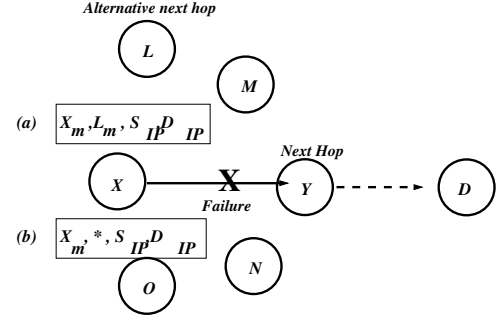


Fig. 2. Two cases when the communication link between X and Y fails.

V. ROUTE MAINTENANCE

A. Refreshing Soft Routing State Using Reverse Traffic

The route refresh mechanism employed by ABP is particularly different from the refresh mechanisms employed by other protocols such as AODV. Each routing entry in ABP routing table has an associated lifetime, which indicates the time duration after which the a non-refreshed routing entry will be deleted. Assume that a node X 's routing table contains a valid routing entry for a destination address D_{IP} with next hop as node Y . The lifetime field of this routing entry is refreshed to a maximum lifetime value of *MAX_ROUTE_LIFETIME* each time a packet *originating* from node D is forwarded by node Y to node X . In other words, X can reaffirm the fact that D is still reachable along the direction of its neighbor Y when it receives a packet from Y that has a *source IP* address of D_{IP} .

The unique aspect of the above route refresh mechanism in ABP is that the **reverse** traffic originating from D is used to refresh the **forward** routing information towards D in the network. This in contrast to other routing protocols such as AODV in which packets *destined* towards D are used to refresh the routing information for D . ABP's refresh mechanism maintains more accurate routing state because only the reverse traffic generated by node D itself can accurately indicate the current forward direction in which D can be reached (assuming the use of bidirectional communication links). On the other hand, with the route refresh mechanism in AODV, just because packets destined towards D are being forwarded using D 's current routing entry does not necessarily mean that the current routing entry is accurate or should be refreshed. This is especially true during node mobility or failures when only the packets originating from D , rather than packets destined towards D , can provide accurate information about D 's reachability.

B. Handling Mobility and Network Failures

Ongoing communication sessions can be adversely affected whenever (1) a node fails, (2) a communication link degrades, or (3) nodes move inducing communication failure. Failure can be detected whenever a node attempts to transmit a unicast packet to one of its neighbors and, in return, does not receive a corresponding link-level acknowledgment. Consider Figure 2 in which there is a failure along a route from S to D between neighbors X and Y . When X detects a failure during unicast

transmission of a packet to its neighbor Y , it first checks every routing entry that contains Y_m as the next hop address. If such a routing entry contains a list of alternative next hops, then the best alternative, as per the link metric used, is assigned as the primary next hop. If there is no alternative next hop, the corresponding routing entry is marked as invalid. For example, in Figure 2, X has four other neighbors L , M , N , and O , besides Y . Of these four, L happens to be the best alternative next hop towards destination D . Upon failure of communication with Y , X first assigns L as the primary next hop and unicasts the packet for D to L . This is shown in case (a). If the communication with L fails as well and there are no more alternative next hops, then X floods the packet to all its neighbors hoping that at least one of them can reach D . This is shown in case (b). The scope of this flood is naturally limited by the packet's TTL value when it is received by X .

Part of the network which is affected by the failure (including node X) needs to quickly re-learn about new routes to the destination D . In case D is still alive and reachable, D is guaranteed to either transmit data packets or dummy IP packets towards all nodes with whom it is communicating (including S). Two possibilities follow. First is that node X lies along the path of these packets being sent by D after the failure. In fact, X may receive packets from D via any of its neighbors L , M , N , O , or Y , and needs to select one as the next hop for D . However X needs to be careful in selecting a stable next hop to D in case the connectivity to D has not yet stabilized before selection. To learn a stable route, X observes all non-duplicate packets from D that arrive via each neighbor for a duration of `MAX_ROUTE_LIFETIME` after the failure. At the end of this duration, X selects the neighbor with the best link/path quality metric as the next hop. For example, one criterion could be to select a neighbor from which X receives the last k consecutive non-duplicate packet sent by D , or a neighbor with highest bandwidth or least signal to noise ratio. The second possibility is that X may not receive any of the packets sent by D after a failure. Consequently, D 's routing entry at node X will expire within `MAX_ROUTE_LIFETIME` duration. This would mean that either D is unreachable or alternate routes have been found that bypass X . In either case, routing entry for D at X would no longer be required.

Note that in other protocols, such as DSR and AODV, if a link failure occurs and there is no alternative route, then a Route Error (RERR) message is sent back to the source. While, this approach has its benefits, it also requires intermediate nodes to recognize and process explicit RERR messages on a per-route basis. Hence there is an inherent tradeoff between the level of protocol complexity and control message processing. In ABP, we resort to the option of limited flooding instead of using explicit control message. Our simulation results indicate that the performance of ABP still remains competitive in spite of this design choice.

C. Soft State Re-learning after Node Reboot

A failure or reboot of a ABP node Y may induce neighboring nodes to re-learn their routes that passed through Y .




Fig. 3. Complete transient loops disappear within provably bounded time.

However, once Y reboots and comes back up, it does not need any re-synchronization with its neighbors. Rather it can immediately start learning and participating in the routing process again. This is in contrast to AODV in which a node is forced to stay silent for a period of time after reboot in order to flush the routing state of its neighbors [21].

VI. DAMPING DUPLICATES WITHOUT A SPANNING TREE

Wired-LAN bridges construct and maintain a network-wide spanning tree to prevent routing loops because the Ethernet header does not carry any hop-count or TTL information [23]. A spanning-tree is neither feasible nor advisable in wireless ad hoc networks due to the higher degree of node mobility and non-determinism in transmissions. In fact, spanning trees are being viewed unfavorably even in wired environments [23], [25]. ABP nodes do not construct or maintain any spanning trees. Instead they damp the propagation of duplicate packets by relying on the Identification and TTL fields carried inside the IP header of every packet. Such a packet can potentially be discarded.

Every ABP node remembers the Identification values of last k packets seen from source node S . A packet is dropped only if its value is among the last k remembered values or is smaller than all of them (after accounting for wrapping around of the 16-bit value). This technique detects duplicates and tolerates re-ordering of packets within a window of k packets. In evaluations, we observe that a window size as small as $k = 3$ is more than sufficient detect and damp all duplicates even in the presence of packet reordering. The worst-case per-node state requirement in an N -node network would be $O(kN)$ (IP Identification field being unique across all connections within a host). Another alternative for duplicate detection with low state overhead could be the use of Bloom Filters [26]. Also note that, unlike the intricate mechanism of destination sequence numbers in AODV, which are used to prevent routing loops altogether and maintain route freshness, the IP Identification field is used in ABP to simply detect and damp duplicate packets.

VII. LIMITING THE IMPACT OF TRANSIENT LOOPS

One side-effect of not relying on a spanning tree is that transient routing loops may arise for some destinations when network topology is changing. We define a *complete* transient

loop for a destination D as one in which each node in the loop has a valid routing entry towards the next hop node in the loop. On the other hand, a *partial* transient loop for a destination D is one in which at least one node in the loop does not have a valid routing entry for D . Consequently, such a node would broadcast the packets destined to D and the next node in the loop picks up the broadcast packet.

Algorithms such as AODV and DSR incorporate intricate mechanisms (based on destination sequence numbers or source routes) to guarantee that *complete* transient loops are never formed. On the other hand, *partial* transient loops are almost impossible to avoid whenever any intermediate node needs to broadcast a packet to an unknown destination. ABP does not attempt to prevent transient loops at all costs. Rather, if and when transient loops are indeed formed, their impact on network performance is highly minimized due to following three factors. (1) *ABP guarantees that complete transient loops last for less than a provably short bounded time interval* (We prove this claim below). (2) Since ABP nodes perform duplicate packet detection (Section VI), no packet will traverse a loop more than once. Rather a packet is discarded once it revisits any node. (3) Successive ABP nodes decrement the TTL field in each packet and discard them when TTL reaches zero. While (2) and (3) guarantee that packets do not proliferate unchecked even when complete or partial transient loops are formed, (1) guarantees that complete transient loops do not last forever.

Theorem 1: Lifetime of a complete transient loop for any destination D is less than or equal to MAX_ROUTE_LIFETIME.

Proof: Consider the complete loop in Figure 3 for destination D with k nodes N_1 to N_k , where the next hop for destination D at node N_i is N_{i+1} and at node N_k is N_1 . By definition, node D itself cannot be part of the complete loop and hence is outside the loop. Furthermore, in order for node N_i to maintain N_{i+1} as its next hop entry for D , N_{i+1} must forward a packet with *source* IP address D_{IP} to N_i within MAX_ROUTE_LIFETIME interval of its learning the next hop. Two cases arise. In the first case, node D does not send a packet through any node in the complete loop for MAX_ROUTE_LIFETIME duration. Thus N_{i+1} cannot forward a packet with D_{IP} as source IP address to N_i . Consequently N_i 's routing entry for D will expire and become invalid, thus breaking the complete loop within MAX_ROUTE_LIFETIME duration. In the second case, node D does send a packet through node N_{i+1} within MAX_ROUTE_LIFETIME interval. Since node D is outside the loop, the packet from D must enter the loop from some external node N_e . Without loss of generality, assume that N_e is a neighbor of node N_{i+1} that we are presently considering. Thus node D 's packet arrives at N_{i+1} from N_e . However, N_e is different from N_{i+2} , which is the current next hop for D at node N_{i+1} (since N_e is outside the loop). Thus, as described in Section V, N_{i+1} will switch to maintenance mode to re-learn its next hop for D by invalidating its routing entry for D via N_{i+2} . This too would break the complete loop within

MAX_ROUTE_LIFETIME interval. ■

It is easy to see from the above argument that a complete loop soon decays into a partial loop in which one or more nodes do not have a valid routing entry towards D and broadcast the packets destined to D . However the partial loop also lasts only until the broadcasting node re-learns a new route to D , which would happen quickly if the broadcasting node lies along the path of data/dummy-IP packets arriving from D . Furthermore, D is *guaranteed* to send either a data or a dummy IP packet within ACTIVITY_INTERVAL, as described in Section V-IV-C.

VIII. PERFORMANCE EVALUATION

In this section we show that, even without explicit control messages, the performance of ABP is highly competitive, if not better, when compared to AODV and DSR protocols. We used the NS2 simulator to evaluate the performance of ABP protocol. Simulations were run for different networks with 50 nodes. The number of connections were varied from 10 to 100, where each connection is set up between a randomly selected node pair. We used the mobility scenarios from the Monarch Project[1] in which the mobility speed of the nodes was varied between from 0 to 20 meters/sec. Nodes move within a rectangular field of 1500m x 300m according to the Random Waypoint Model. Each simulation is run for 900-simulated seconds. Pause times were varied from 0s (Highly mobile) to 900s (static). We used both TCP-based connections, (for bidirectional communication), and UDP-based Constant Bit Rate (CBR) connections (for unidirectional communication). The maximum TCP window was set to 20 packets, with maximum packet size set to 1000 bytes. The UDP CBR sources sent traffic at 4 packets/sec with 64 byte packets. Each data point is averaged over five simulation runs with a different random seed to vary the connection establishment pattern.

ABP does not use any explicit control messages, ARP requests/replies or Hello messages. In the presence of unidirectional data traffic, ABP nodes generate implicit data-like control messages called *dummy* IP packet. In these simulations, we count the dummy IP packets as the control packets for ABP. The MAX_ROUTE_LIFETIME is set to 5 seconds and the dummy IP packets are generated by the ABP layer at a node after 4 seconds of silence on the part of the network application. For AODV and DSR, we count the Route Request (RREQ), Route Reply (RREP), Route Error (RERR), ARP requests/replies, and Hello messages as the control overhead. For DSR, we *do not* count the source routing information carried in packet headers towards the control overhead. The control packets are counted on a hop-to-hop basis rather than on end-to-end basis. Thus, if a control packet is forwarded or broadcasted k times, then it is counted as k control messages.

For fair comparison, in AODV simulations we enabled the Local Repair optimization (which attempts to repair a failed link locally before returning a RREQ message to source) and the Link Layer Detection optimization (which suppresses the transmission of "Hello" messages). Similarly, the simulation of DSR included the optimization of piggybacking RERR and

RREP messages onto RREQ messages and the optimization of populating route cache by listening in promiscuous mode. ABP and AODV simulations do not use promiscuous mode.

We use three metrics : (1) *R/S* is the ratio of the number of packets received by the destination to the number of packets sent by the source. (2) *C/R* is the ratio of number of control packets to the number of data packets received by destination. It signifies the number of control packets sent for each data packet received. (3) *Number of control packets* represents the total number of control messages transmitted during the entire simulation. We note that, among the three metrics, *R/S* represents the most accurate measure of protocol performance. This is because *R/S* incorporates the impact of both control messages (both non-piggybacked and piggybacked) as well as data packets (both unicast and broadcast) on the overall protocol performance.

A. Performance with TCP Traffic

R/S: Figures 4, 5, and 6 show that DSR performs best in terms of *R/S* values, followed by ABP and AODV. In Figure 4, as the pause time increases, the *R/S* value approaches 1 (no packet loss). In Figure 5, as the number of connections increases, channel contention and packet losses increase resulting in a drop in *R/S*. As in the previous case, DSR performs best, followed by ABP and finally AODV. Figure 6 shows a universal decrease in *R/S* with increase in mobility speed. The difference in *R/S* between DSR and ABP is less than 1%.

It can be seen from these figures that reducing the protocol complexity in ABP does not automatically translate into better delivery ratio (*R/S*) when compared against DSR. This is because reduction in control message overhead is offset to some extent by increase in both unicast and broadcast data traffic, and this reflects in the final *R/S* numbers. However, ABP's delivery ratio is still competitive with AODV and DSR.

C/R: Figures 7, 8, and 9 show the impact of various parameters on *C/R* values. In all the three figures, ABP maintains a low value of *C/R* around 0.2 because it exploits the bidirectional nature of TCP traffic to maintain routing state without exchanging excessive control messages. DSR and AODV show higher value of *C/R* for smaller pause times than for larger pause times. These figures show that larger pause times, smaller number of connections and lower mobility speed, all contribute to less dynamic network and smaller *C/R*. ABP experiences little variation in *C/R* whereas variation in *C/R* is large for both AODV and DSR. It is notable that DSR, which relies upon source routing, performs better than ABP by a difference in *C/R* of 0.05 for large pause times. Figure 9 shows that increasing the mobility speed of nodes results in greater increase in *C/R* for DSR than for ABP and AODV. One possible reason is that DSR employs source routing as opposed to ABP and AODV that use backward learning. Consequently, DSR is unable to benefit from the forward routing information carried in the reverse traffic.

Number of Control Packets: Figure 10 shows that number of control messages in ABP is much lower than AODV and DSR. For 0s pause time, ABP generates around 20K control

packets, AODV generates around 90K control packets and DSR generates around 120K control packets. Thus, in comparison, ABP generates 4.5 times fewer control packets than AODV and 6 times fewer control packets than DSR. In Figure 11, we vary the number of TCP connections from 10 to 150, for a fixed 60s pause time and fixed mobility speed of 20m/s. One can see that number of control packets generated by ABP remains much lower than AODV and DSR. As the number of connections increases, AODV generates more control packets than DSR. For 150 connections, AODV generates around 220K control packets, DSR generates around 125K control packets, while ABP generates around 25K control packets. As with *C/R*, Figure 12 shows that increasing the mobility speed of nodes results in greater increase in number of control packets for DSR than for ABP and AODV.

B. Performance with Constant Bit Rate (CBR) UDP Traffic

In this section, due to space constraints, we omit the mobility speed variation which shows similar trends as pause time and the UDP connection count variation.

R/S: In Figure 13, ABP achieves an *R/S* smaller than both AODV and DSR for smaller pause times. As pause time increases, ABP's performance approaches that of AODV. In Figure 14, ABP and DSR achieve almost the same *R/S* which is smaller than AODV. When compared against results for TCP traffic, the above results show that, as expected, ABP performs better in the presence of bidirectional traffic, than in the presence unidirectional traffic.

C/R: Figure 15 shows that AODV transmits almost 1.5 control packets for each received data packet. Although, DSR has approximately 0.8 *C/R* at 0s pause time, it reduces as the pause time increases. ABP has a *C/R* less than 0.2 at any given point of time. Figure 16 shows that, as the number of connections increases, the *C/R* value drops for ABP. With more number of connections, dummy IP packets are suppressed more often leading to lower *C/R* value. On the other hand, *C/R* value increases for both DSR and AODV.

Control Packets: In Figure 17, AODV generates around 140K control packets for 0s pause time, whereas DSR generates around 75K control packets and ABP generates around 12K control packets. In Figure 18 we can see that, although all routing protocols use fewer control messages for smaller number of connections, the main difference arises when the number of connections increases. For 150 connections AODV uses around 320K control packets and DSR uses around 180K control packets, while ABP uses less than 20K control packets.

IX. RELATED WORK

Numerous ad hoc routing algorithms have been proposed in literature, all of which rely upon the use of extensive control infrastructure for loop prevention and/or explicit control messages that requires special processing in the network interior. Here we review the ones most relevant to this paper. Existing ad hoc routing protocols can be classified as either *table-driven* or *on-demand* protocols. DSDV [6] is a table-driven protocol based on the distance vector algorithm and provides guarantees




Fig. 4. R/S vs Pause Time for 50 nodes, 30 TCP connections, 20m/s.




Fig. 5. R/S vs TCP connections for 50 nodes, 60s pause time, 20m/s.




Fig. 6. R/S vs Mobility speed for 50 nodes, 30 TCP connections, 60s pause time.




Fig. 7. C/R vs Pause Time for 50 nodes, 30 TCP connections, 20m/s.




Fig. 8. C/R vs TCP Connections for 50 nodes, 60s pause time, 20m/s.




Fig. 9. C/R vs Mobility speed for 50 nodes, 30 TCP connections, 60s pause time.




Fig. 10. Control packets vs Pause Time for 50 nodes, 30 TCP connections, 20m/s.




Fig. 11. Control packets vs TCP Connections for 50 nodes, 60s pause time, 20m/s.




Fig. 12. Control pkts vs Mobility speed. 50 nodes, 30 TCP connections, 60s pause time.




Fig. 13. R/S vs Pause Time for 50 nodes, 30 UDP connections, 20m/s.




Fig. 14. R/S vs Number of UDP connections for 50 nodes, 60s pause time, 20m/s.




Fig. 15. C/R vs Pause Time for 50 nodes, 30 UDP connections, 20m/s.




Fig. 16. C/R vs UDP Connections for 50 nodes, 60s pause time, 20m/s.




Fig. 17. Control pkts vs Pause Time for 50 nodes, 30 UDP connections, 20m/s.




Fig. 18. Control pkts vs UDP Connections for 50 nodes, 60s pause time, 20m/s.

on loop freedom by use of destination sequence numbers. DSDV uses explicit control messages to proactively exchange link-state information. A number of other table-driven protocols exist in literature [18], [9], [24], [20], all of which inherently require the exchange of explicit control messages to proactively maintain routes. *Ad hoc On-Demand Distance Vector (AODV)*[22] routing protocol is an on-demand version of distance vector algorithm that establishes a route only when required. Like DSDV, AODV guarantees loop freedom through use of destination sequence numbers. AODV relies upon explicit control messages such as RREQ, RREP, and RERR messages to construct and maintain routes. *Dynamic Source Routing (DSR)*[12] is an on-demand routing protocol that uses source routing to guarantee loop freedom. DSR uses explicit control messages, such as RREQ and RREP packets, to accumulate the address of each intermediate node during the route discovery phase. The accumulated source route is then appended to each packet's header before transmission. LQSR [7] is a variant of DSR that supports link quality metrics. A number of other on-demand protocols have been proposed in literature, [19], [10], [15], [4], [17], [3] to name a few, all of which use explicit control messages. One class of ad hoc routing protocols [5], [11], [8] use the hierarchical structure of addressing to make routing decisions. Another class of routing protocols [14], [16], [2], [13] use location information to route data packets.

To the best of our knowledge, all the above protocols use explicit control messages and some form of control infrastructure to construct and maintain loop-free routes in the face of node mobility and communication failures. Some of the above mentioned protocols try to minimize the number of control messages by piggybacking explicit control information with data or other control packets. For instance, DSR can piggyback RERR and RREP messages over RREQ messages and includes the complete source routing information within every packet. In such cases, intermediate nodes still need to understand and process these protocol-specific explicit control headers and need additional control logic to maintain loop-free routes. In contrast, ABP does not rely upon the use of any explicit control infrastructure or control information, whether piggybacked or otherwise, and yet achieves competitive network performance.

X. CONCLUSION

In this paper, we have proposed a new wireless ad hoc routing protocol, called *Ad hoc Bridging Protocol (ABP)*, which is inspired from the self-learning plug-and-play nature of transparent bridges in wired LANs. The distinguishing feature of ABP is that it does not use any explicit control messages and does not rely upon any extensive control infrastructure for loop prevention (such as spanning trees, destination sequence numbers, or source routing mechanisms). ABP is particularly suited for routing bidirectional communication sessions with minimal control overhead. To handle unidirectional communication, ABP makes minimal use of implicit (data-like) control messages that require no special processing at intermediate nodes. We prove that transient loops, if and when they form,

do not last longer than a bounded short duration and do not negatively impact network performance. We also demonstrate that, even without explicit control messages, ABP's performance is quite competitive compared to AODV and DSR. We are currently investigating issues of security and disruption in ABP, which are fundamentally the same as in any protocol that relies upon backward learning principle.

REFERENCES

- [1] The monarch project. <http://www.monarch.cs.rice.edu/>.
- [2] S. Basagni, I. Chlamtac, V. Syrotiuk, and B. Woodward. A distance routing effect algorithm for mobility. In *MobiCom*, 1998.
- [3] S. Biswas and R. Morris. Exor: opportunistic multi-hop routing for wireless networks. In *Proc. of SIGCOMM*, New York, NY, USA, 2005.
- [4] I. Chakeres, E. Royer, and C. Perkins. Dynamic MANET on-demand routing protocol. IETF Internet Draft, October 2005.
- [5] T. Chen and M. Gerla. Global state routing: A new routing scheme for ad-hoc wireless networks. In *ICC'98, Atlanta, GA*, June 1998.
- [6] C. Perkins and P. Bhagwat. Highly dynamic destination-sequenced distance-vector routing (DSDV) for mobile computers. In *ACM SIGCOMM'94*, pages 234–244, 1994.
- [7] R. Draves, J. Padhye, and B. Zill. Routing in multi-radio, multi-hop wireless mesh networks. In *Proc. of MobiCom*, New York, NY, USA, 2004.
- [8] J. Eriksson, M. Faloutsos, and S. Krishnamurthy. Scalable ad hoc routing: The case for dynamic addressing. In *INFOCOM*, 2004.
- [9] J. J. Garcia-Luna-Aceves and M. Spohn. Source-tree routing in wireless networks. In *ICNP*, 1999.
- [10] Z. Haas, M. Pearlman, and P. Samar. The interzone routing protocol for ad hoc networks. IETF Internet draft, July 2002.
- [11] A. Iwata, C. Chiang, G. Pei, M. Gerla, and T. Chen. Scalable routing strategies for ad-hoc wireless networks. *IEEE JSAC*, 17(8):1369–1379, Aug. 1999.
- [12] D. B. Johnson and D. A. Maltz. Dynamic source routing in ad hoc wireless networks. In *Mobile Computing*, volume 353. Kluwer Academic Publishers, 1996.
- [13] B. Karp and H.T.Kung. Greedy perimeter stateless routing for wireless networks. *MobiComm 2000*, pages 1–8, 2000.
- [14] Y. Ko and N. Vaidya. Location-aided routing in mobile ad hoc networks. In *MOBICOM'98*, October 1998.
- [15] Y. Lee and G. Riley. Dynamic NIX-Vector routing for mobile ad hoc networks. In *Wireless Commun. and Networking Conf.*, 2005.
- [16] J. Li, J. Jannotti, D. De Couto, D. Karger, and R. Morris. A scalable location service for geographic ad-hoc routing. In *Mobicom*, 2000.
- [17] B. Manoj, R. Ananthapadmanabha, and C. Murthy. Link life based routing protocol for ad hoc wireless networks. In *ICC*, Oct. 2001.
- [18] S. Murthy and J. Garcia-Luna-Aceves. A routing protocol for packet radio networks. In *Mobile Computing and Networking*, 1995.
- [19] V. Park and M. Corson. A highly adaptive distributed routing algorithm for mobile wireless networks. In *INFOCOM*, 1997.
- [20] T. Parker and K. Langendoen. Guesswork: Robust routing in an uncertain world. In *Mobile Ad-hoc and Sensor Systems*, 2005.
- [21] C. Perkins, E. Belding-Royer, and S. Das. Ad-hoc on demand distance vector routing. RFC 3561, July 2003.
- [22] C. Perkins and E. M. Royer. Ad-hoc on-demand distance vector routing. *IEEE Workshop on Mobile Comp Sys and Applications*, 1999.
- [23] R. Perlman. Rbridges: Transparent routing. In *Proc. of Infocom 2004*, Hong Kong, China, June 2004.
- [24] C. Santivanez and R. Ramanathan. Hazy sighted link state routing protocol. BBN Technical Memorandum No. 1301, Aug. 31 2001.
- [25] S. Sharma, K. Gopalan, S. Nanda, and T. Chiueh. Viking: A multi-spanning-tree ethernet architecture for metropolitan area and cluster networks. In *Proc. of Infocom 2004*, Hong Kong, China, June 2004.
- [26] A. C. Snoeren. Hash-based ip traceback. In *Proc. of the SIGCOMM*, 2001.